


TiKV 高性能 Rust 开发实践

Wenxuan @  PingCAP  KV





施闻轩 · TiKV Engineer


TiKV Optimization & Development

TiKV Coprocessor SIG Tech Lead


Sysbench

2.1 → 3.0

Sysbench - Update Non-Index




3.0 → 4.0 (not released)




TPC-H

2.1 → 3.0

TPC-H 50



3.0 → 4.0 (not released)



Performance boost comes from...

1. New Feature / Algorithm / Data Structure

- v3.0: Multi-thread Raft
- v3.0: Batch gRPC messages
- v3.0: Titan Engine
- v3.0: Vectorization
- v4.0: Will be officially announced in future

2. Engineering Improvement / Constant Optimization

Rust Claims...

Rust

A language empowering everyone
to build reliable and efficient software.

GET STARTED

[Version 1.39.0](#)

Why Rust?

Performance

Rust is blazingly fast and memory-efficient: with no runtime or garbage collector, it can power performance-critical services, run on embedded devices, and easily integrate with other languages.

Reliability

Rust's rich type system and ownership model guarantee memory-safety and thread-safety — and enable you to eliminate many classes of bugs at compile-time.

Productivity

Rust has great documentation, a friendly compiler with useful error messages, and top-notch tooling — an integrated package manager and build tool, smart multi-editor support with auto-completion and type inspections, an auto-formatter, and more.

Rust Claims...

Rust

GET STARTED

Version 1.39.0

doesn't mean that you can
write efficient code unthinkingly

Rust is blazingly fast and memory-efficient: with no runtime or garbage collector, it can power performance-critical services, run on embedded devices, and easily integrate with other languages.

Rust's rich type system and ownership model guarantee memory-safety and thread-safety — and enable you to eliminate many classes of bugs at compile-time.

Rust has great documentation, a friendly compiler with useful error messages, and top-notch tooling — an integrated package manager and build tool, smart multi-editor support with auto-completion and type inspections, an auto-formatter, and more.

Find Hotspots



Why?

Amdahl's law


$$S_{\text{latency}}(s) = \frac{1}{(1 - p) + \frac{p}{s}}$$

perf + FlameGraph

<https://github.com/brendangregg/FlameGraph>

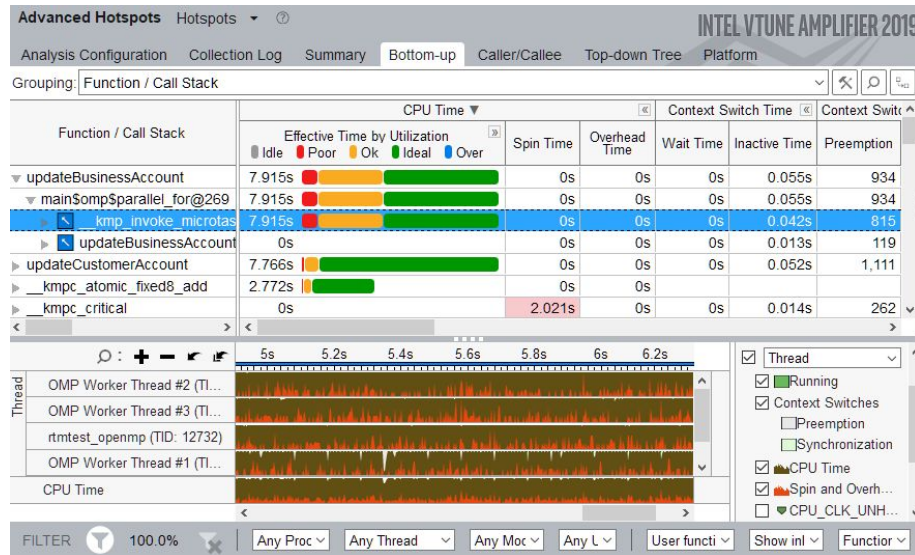
```
$ perf record ...
$ perf script |
  ./stackcollapse-perf.pl |
  ./flamegraph.pl
```

Even very useful in production!



Intel VTune

- It is free now
- Easy to use GUI
- Top-down Microarchitecture Analysis



Low-hanging Fruits



Replace Memory Allocators

Use efficient memory allocators like [jemalloc](#):

```
#[global_allocator]  
static GLOBAL: Jemalloc = Jemalloc;
```

Improve ~25% for TPC-H Q1.

Link Time Optimization

LTO can apply cross-crate and even cross-language optimizations!

```
[profile.release]
```

```
lto = true
```

Improve ~15% for TPC-H Q1.

Set Target CPU

If not set (default):

- Portable release (e.g. x86_64)
- Modern instructions like AVX is not used

(Almost) Drop-in Replacement

- **Hash Collections** like `std::collections::HashMap` → [hashbrown](#) (*already in std*)
 - 2.6x faster insert
- **Concurrent utilities** like `std::sync::mpsc` → [crossbeam](#)
 - It is even faster than Go channel
- **Synchronization primitives** like `std::sync::Mutex` → [parking_lot](#)
 - 5x faster Mutex

Micro Optimizations



Compiler Can Be Stupid

```
1 pub fn foo(v: Vec<u64>) -> Vec<u64> {  
2     v.into_iter().map(|v| v).collect()  
3 }  
4 |
```



```
34 example::foo:  
35     push    rbp  
36     push    r15  
37     push    r14  
38     push    r13  
39     push    r12  
40     push    rbx  
41     sub     rsp, 72  
42     mov     r12, rdi  
43     mov     rbx, qword ptr [rsi]  
44     mov     rax, qword ptr [rsi + 8]  
45     mov     r15, qword ptr [rsi + 16]  
46     lea     r14, [rbx + 8*r15]  
47     mov     qword ptr [rsp], 8  
48     xorps   xmm0, xmm0  
49     movups  xmmword ptr [rsp + 8], xmm0  
50     mov     qword ptr [rsp + 40], rbx  
51     mov     qword ptr [rsp + 48], rax  
52     mov     qword ptr [rsp + 56], rbx  
53     mov     qword ptr [rsp + 64], r14  
54     test    r15, r15  
55     je      .LBB2_1  
56     mov     qword ptr [rsp + 24], rax  
57     mov     qword ptr [rsp + 32], r12  
58     lea     r12, [8*r15]  
59     sar     r12, 3  
60     mov     ecx, 8  
61     xor     r13d, r13d  
62     mov     rax, r12  
63     mul     rcx  
64     mov     rbp, rax  
65     setno   al  
66     jo      .LBB2_5  
67     mov     r13b, al  
68     shl     r13, 3  
69     mov     rdi, rbp
```

Less Allocation, Less Copy

Pre-allocate collections when possible:

- `Vec::with_capacity()`

Less Allocation, Less Copy

Use memory pool techniques to reduce allocation:

- Dynamic collections over pre-allocated memory: [bumpalo](#)
- Dynamically allocate same object: [TypedArena](#)
- 1d `Vec<Vec<T>>`: [nested](#)

Less Allocation, Less Copy

Prefer stack allocation instead of heap allocation:

- Static array allocates on stack: `[T; N]`
- `Box<T>` allocates on heap
- `Vec<T>` allocates data on heap (obviously)
- `SmallVec` can be very useful
 - allocates N bytes on stack but is also growable (on heap for `len > N`).

Less Allocation, Less Copy

Move less data

- `Box<[T]>` is smaller than `Vec<T>`
- Wrap *big* and *frequently moved* structures with `Box<T>` to keep it small

Less Allocation, Less Copy

Error size matters:

- `Result<T, E>` is an enum: `sizeof(Result<T, E>) == max(sizeof(T), sizeof(E))`
- The size of `E` is the minimum size of your `Result`.
- Try `Result` is not zero cost sometimes!

TiKV TPC-H improved by 14% by reducing errors from 200 bytes to 8 bytes:

- Reduced 3×200 bytes memory copy each KV iterate
- `Result<T, Error> → Result<T, Box<Error>>`

Less Allocation, Less Copy

Use reference where possible:

- Annotate the structure with a lifetime parameter to contain references.

```
pub struct WriteRef<'a> {  
    pub write_type: WriteType,  
    pub start_ts: Timestamp,  
    pub short_value: Option<&'a [u8]>,  
}
```


Less Allocation, Less Copy

Use reference where possible:

- Use `Arc<T>` instead of a raw reference when lifetime is hard to write.

Less Allocation, Less Copy

Use reference where possible:

- Use DSTs when you only want to wrap a DST reference (like `&[u8]`).

Example: Need `&[u8]` internally, but encapsulate with a type to provide extra feature.

```
struct Key(Vec<u8>);  
fn get(k: &Key);
```

✗ Not zero cost

```
struct KeyRef<'a>(&'a [u8]);  
fn get(k: KeyRef<'_>);
```

✓ Zero cost, but interface doesn't
accept reference any more

```
struct KeyRef([u8]);  
fn get(k: &KeyRef);
```

✓ Best: DST wrapper

Less Allocation, Less Copy

Use reference where possible:

- Unbounded lifetime / raw pointer:

```
unsafe fn erase_lifetime<'a, T: ?Sized>(v: &T) -> &'a T {  
    &*(v as *const T)  
}
```

- Useful to simulate a '**Self**' lifetime.
- Hate **unsafe**{ } ? [owning_ref](#) or [rental](#) can help.
 - *However sometimes “unsafe” but **clear & simple** code is better.*

Prefer Static Dispatch

```
fn foo(callback:  
    Box<dyn FnOnce()>  
);
```

✗ Box & Dynamic Dispatch

```
fn foo<F: FnOnce()>  
    (callback: F);
```

✓ Static Dispatch

```
fn foo(callback:  
    impl FnOnce()  
);
```

✓ Syntax Sugar

Use `Box<dyn T>` only when you want **different base types** in one type:

- Return **different base types** in one function: `fn foo() -> Box<dyn T>`
- Store **different base types** in one container: `Vec<Box<dyn T>>`

Reduce Branches

- To reduce instructions
- To reduce misprediction

Example: Memory-comparable encoding bytes (padding = 4).

F	o	o
---	---	---

Source

```
fn memcmp_encode(bytes: &[u8], is_desc: bool);
```

F	o	o	\1	\FE
---	---	---	----	-----

Asecnding

× Branch

F	o	o	\FE	\1
---	---	---	-----	----

Descending

Reduce Branches

- To reduce instructions
- To reduce misprediction

Example: Memory-comparable encoding bytes (padding = 4).

F	o	o
---	---	---

Source

F	o	o	\1	\FE
---	---	---	----	-----

Asecnding

F	o	o	\FE	\1
---	---	---	-----	----

Descending

```
fn memcmp_encode_asc(bytes: &[u8]);
```

```
fn memcmp_encode_desc(bytes: &[u8]);
```

✓ No branch, some duplicate code

Reduce Branches

- To reduce instructions
- To reduce misprediction

Example: Memory-comparable encoding bytes (padding = 4).

F	o	o
---	---	---

Source

F	o	o	\1	\FE
---	---	---	----	-----

Asecnding

F	o	o	\FE	\1
---	---	---	-----	----

Descending

```
trait Mode { const FLIP_BITS: bool; }

struct ModeAsc;
impl Mode for ModeAsc {
    const FLIP_BITS: bool = false;
}
struct ModeDesc;
impl Mode for ModeDesc {
    const FLIP_BITS: bool = true;
}

fn memcmp_encode(bytes: &[u8], mode: impl Mode);
```

Reduce Branches

- To reduce instructions
- To reduce misprediction

Example: Memory-comparable encoding bytes (padding = 4).

F	o	o
---	---	---

Source

F	o	o	\1	\FE
---	---	---	----	-----

Asecnding

F	o	o	\FE	\1
---	---	---	-----	----

Descending

```
fn memcmp_encode<const T: bool>  
  (bytes: &[u8], is_desc: T);
```

✓ No branch (const generics) ^{NIGHTLY}

Useful Friends



Micro Benchmark: cargo bench

```
#[bench]
fn bench_foo(b: &mut Bencher) {
    b.iter(|| foo());
}
```

Run benchmark:

```
$ cargo bench
```

Micro Benchmark: cargo bench

Remember to add `black_box`:

```
#[bench]
fn bench_add(b: &mut Bencher) {
    b.iter(|| add(1, 2));
}
```

✗ Constants are likely to be optimized away

```
#[bench]
fn bench_add(b: &mut Bencher) {
    b.iter(|| add(
        black_box(1),
        black_box(2),
    ));
}
```

✓ Wrap constants with `black_box`

Micro Benchmark: criterion

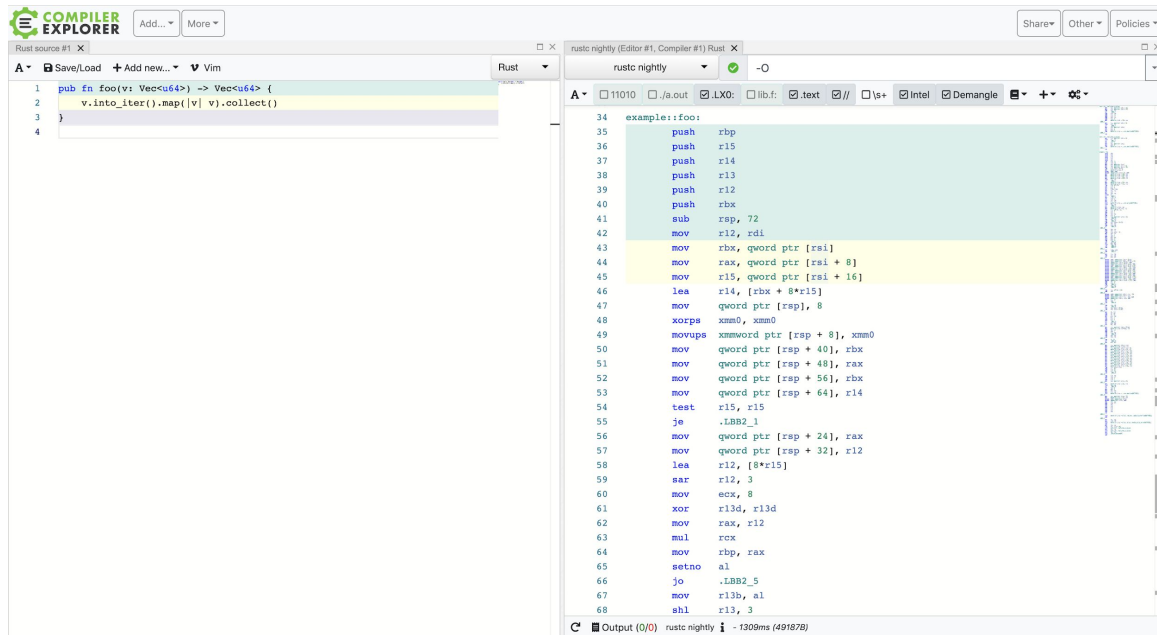
[criterion](#) is a replacement for cargo bench.

Features:

- Statistics-driven, avoid measurement noise
- Support different measurements (e.g. number of instructions instead of wall time)
- Generate charts
- Support parameterized benchmark

Compiler Explorer

<https://godbolt.org/>



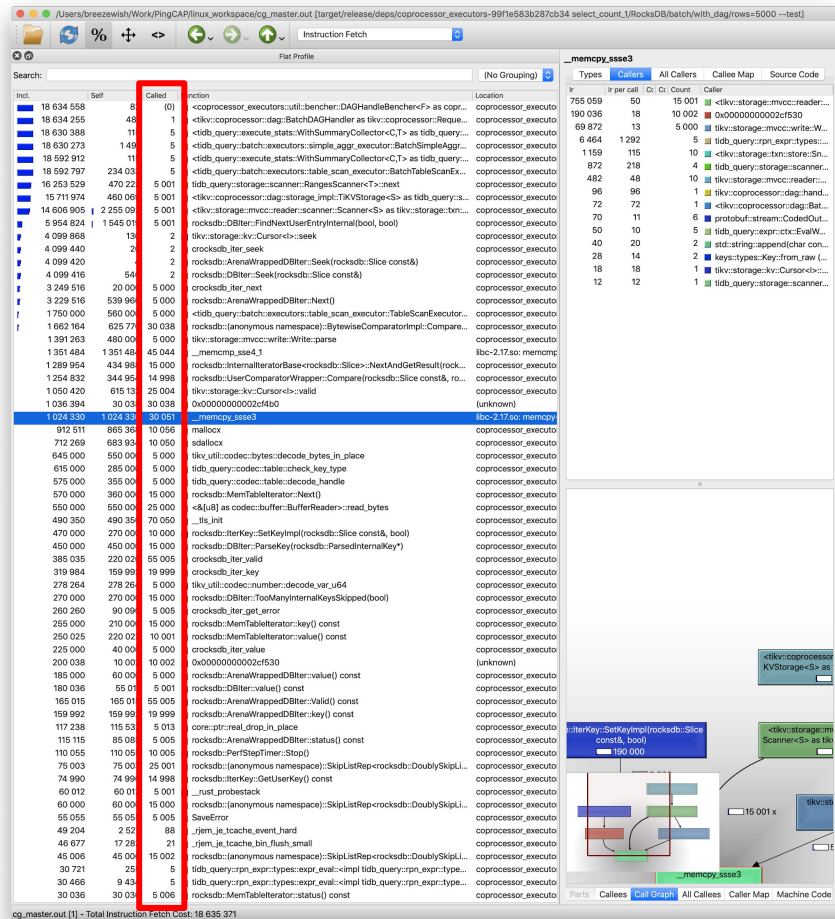
The screenshot displays the Compiler Explorer (Godbolt) interface. The top bar features the 'COMPILER EXPLORER' logo, a search bar, and buttons for 'Add...', 'More', 'Share', 'Other', and 'Policies'. The main window is divided into two panes. The left pane, titled 'Rust source #1', shows the following Rust code:

```
1 pub fn foo(v: Vec<u64>) -> Vec<u64> {  
2     v.into_iter().map(|v| v).collect()  
3 }  
4
```

The right pane, titled 'rustc nightly (Editor #1, Compiler #1) Rust', shows the generated assembly code for the 'foo' function. The assembly is in x86_64 format and includes instructions such as 'push', 'mov', 'leaq', 'xor', 'mul', 'setno', and 'shl'. The assembly is color-coded to match the Rust code in the left pane. The bottom status bar indicates the compiler is 'rustc nightly' and the output is 'Output (0/0)'.

Callgrind + KCacheGrind

- Heavyweight
- Not a sample based profiler
- Run in Valgrind
- **Precise** function calls
- You can use the [callgrind crate](#) to control start / stop (to skip counting unwanted part)



Want to Contribute?




TiKV Community

- [TiKV Help Wanted Issues](#)
- [TiKV PCP Issues](#)
- TiDB Hackathon

PCP Season 1



- [TiDB Performance Challenge Program](#)
- Focus on *performance improvement*
- Season 1: 2019.11.04 - 2020.02.04
- You will be mentored
- You can be rewarded



Special Interest Group

- Focus on specific module.
- Promote to SIG Reviewer / SIG Committer.
- You can receive SIG credit and PCP credit at the same time!
- [TiKV Coprocessor SIG](#)
- [TiKV Engine SIG](#)
- [TiDB Expression SIG](#)

Thank You !



Infra Meetup No. 120 活动交流
群



该二维码 7 天内 (12月13日前) 有效, 重新进入将更新