

Horoscope and TiDB Query Optimizer

Jian Zhang



About Me

- Zhang Jian (张建), TiDB Product and Tech Manager
- zz-jason on GitHub
- Focused on:
 - Query Optimization
 - Distrubited Computation
 - Scheduling
- Email: zhangjian@pingcap.com



Part I - Background



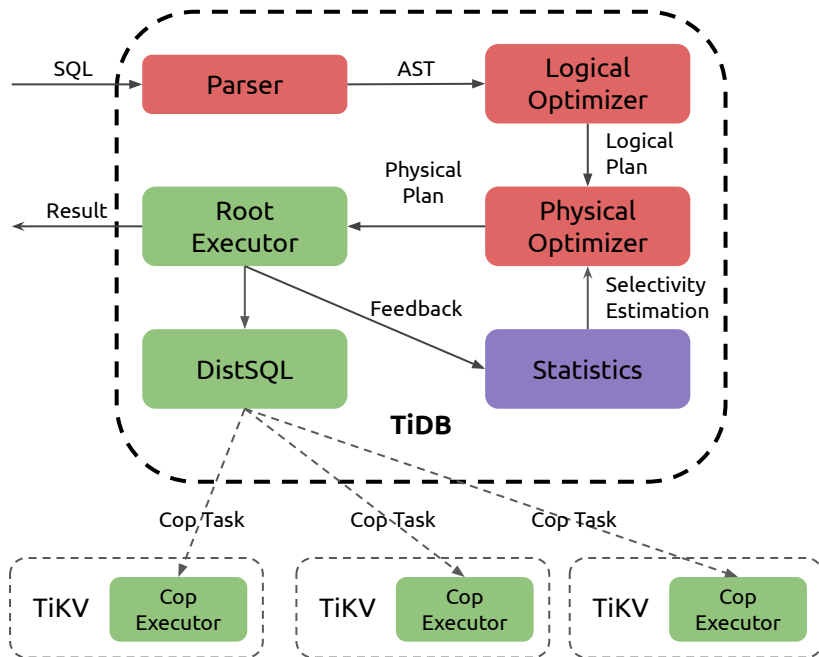
Brief Introduction to TiDB Query Optimizer

Logical Optimization: Equal, Beneficial

- Column Pruning
- Partition Pruning
- Group By Elimination
- etc.

Physical Optimization: Dynamic Programming

- Which index to choose?
- Hash Join, Merge Join, or Index Join?
- etc.



Questions About the Optimizer

How good is the query optimizer?

- What's the percentage that the optimizer cannot choose the best index?
- Is the plan generated for a query the best one?
- How to measure the estimation errors?
- Is the optimizer better than the old version?
- etc.



Part II - Horoscope



Optimizer Test in Other DBMSs

Highly Recommend: <https://github.com/zhangysh1995/awesome-database-testing>



Basic Idea

Inspired by ***OptMark: A Toolkit for Benchmarking Query Optimizers*** and ***Counting, Enumerating, and Sampling of Execution Plans in a Cost-Based Query Optimizer***

Let's enumerate all the query plans, execute them

<https://github.com/chaos-mesh/horoscope>



How to Enumerate All the Execution Plans

```
TiDB(root@127.0.0.1:test) > explain select /*+ nth_plan(1) */ * from t where a = 1 and b > 0 and b < 10;
```

id	estRows	task	access object	operator info
TableReader_7	0.25	root		data:Selection_6
└─Selection_6	0.25	cop[tikv]		eq(hehe.t.a, 1), gt(hehe.t.b, 0), lt(hehe.t.b, 10)
└─┬─TableFullScan_5	10000.00	cop[tikv]	table:t	keep order:false, stats:pseudo

```
3 rows in set (0.00 sec)
```

```
TiDB(root@127.0.0.1:test) > explain select /*+ nth_plan(2) */ * from t where a = 1 and b > 0 and b < 10;
```

id	estRows	task	access object	operator info
IndexLookUp_11	0.25	root		
├─IndexRangeScan_8(Build)	10.00	cop[tikv]	table:t, index:idx_a(a)	range:[1,1], keep order:false, stats:pseudo
└─Selection_10(Probe)	0.25	cop[tikv]		gt(hehe.t.b, 0), lt(hehe.t.b, 10)
└─┬─TableRowIDScan_9	10.00	cop[tikv]	table:t	keep order:false, stats:pseudo

```
4 rows in set (0.00 sec)
```

The First Test Report on Optimizer Effectiveness

Optimizer Effectiveness on TPC-H(SF=10) Test Report, 2020-07-08. ([tidb/issues/18431](https://github.com/pingcap/tidb/issues/18431))

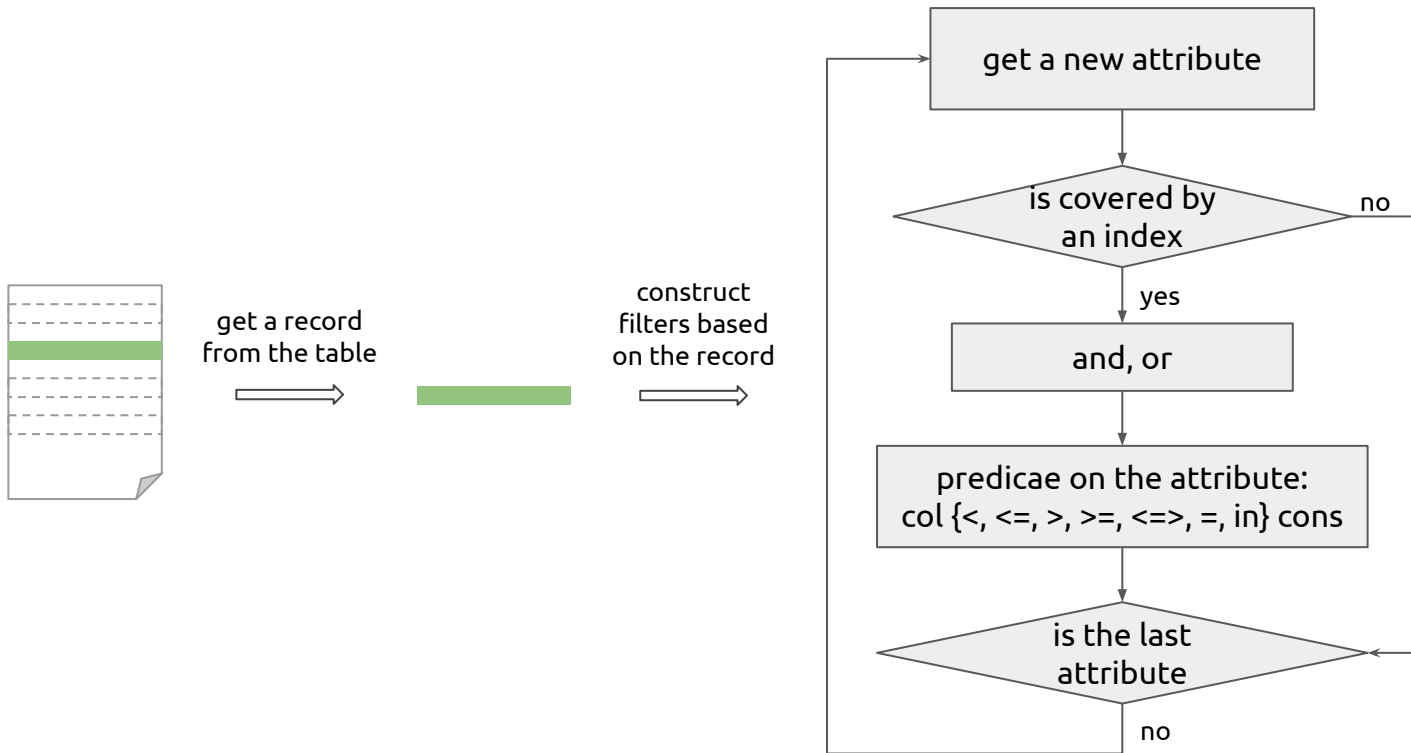
ID	#PLAN SPACE	DEFAULT EXECUTION TIME	BEST PLAN EXECUTION TIME	EFFECTIVENESS	BETTER OPTIMAL PLANS
q1	2	15947.0ms ±107%	15947.0ms ±107%	100.0%	
q2	78	2600.2ms ± 2%	2600.2ms ± 2%	100.0%	
q3	11	11108.8ms ± 3%	6868.8ms ± 4%	72.7%	#6(71.1%),#10(68.4%),#11(61.8%)
q4	5	3670.5ms ± 2%	3670.5ms ± 2%	100.0%	
q5	15	7889.8ms ± 3%	7889.8ms ± 3%	100.0%	
q6	4	4375.2ms ± 3%	4167.8ms ± 5%	50.0%	#1(95.3%),#4(97.6%)
q7	18	6506.8ms ± 8%	6506.8ms ± 8%	100.0%	
q8	23	7167.5ms ± 3%	7167.5ms ± 3%	100.0%	
q9	18	25325.5ms ± 5%	21859.5ms ± 4%	77.8%	#11(93.3%),#12(93.2%),#17(86.3%),#18(89.6%)
q10	10	4941.0ms ± 4%	4556.8ms ± 5%	40.0%	#1(92.2%),#2(93.4%),#3(93.3%),#4(94.4%),#5(92.6%),#8(94.7%)
q11	6	2468.2ms ± 4%	2468.2ms ± 4%	100.0%	
q12	9	5707.8ms ± 4%	5707.8ms ± 4%	100.0%	
q13	2	6343.8ms ± 1%	6343.8ms ± 1%	100.0%	
q14	10	5090.2ms ± 4%	4823.2ms ± 7%	80.0%	#4(96.4%),#5(94.8%)
q16	13	2740.8ms ±10%	2740.8ms ±10%	100.0%	
q17	13	18695.5ms ± 1%	17249.0ms ± 3%	69.2%	#8(92.3%),#9(99.2%),#11(93.4%),#13(92.6%)
q18	23	30369.0ms ± 3%	23088.2ms ± 7%	56.5%	#13(84.4%),#14(91.8%),#15(94.6%),#16(85.1%),#17(91.7%),#18(96.1%),#19(76.0%),#20(84.9%),#21(83.9%),#22(94.2%)
q19	10	6423.0ms ± 1%	6423.0ms ± 1%	100.0%	
q20	51	4943.8ms ± 6%	4943.8ms ± 6%	100.0%	
q21	16	10854.0ms ± 3%	10619.2ms ± 2%	93.8%	#8(97.8%)
q22	1	4122.2ms ± 8%	4122.2ms ± 8%	100.0%	

What Defines a Test Case

- Dataset: There're lots of real-world dataset on the internet.
- Schema:
 - Which type to use for an attribute?
 - What index to construct for a table?
- Query:
 - How to generate queries based on the schema?
 - How to generate queries which cover most of the application use cases?



How to Generate Queries



How to Generate Test Cases

- Generate add-indexes DDLs

```
$ horo index gen
```

- Apply add-indexes DDLS

```
$ horo index add
```

- Generate queries

```
$ horo gen
```

```
SELECT *  
FROM info_type,  
      kind_type,  
      link_type  
WHERE ((info_type.id <=> 49  
        OR info_type.id < 49)  
        AND (info_type.info <=> 'LD spaciality'  
              OR info_type.info > 'LD spaciality'))  
AND ((kind_type.id <=> 6  
       OR kind_type.id < 6)  
      OR (kind_type.kind <=> 'video game'  
          OR kind_type.kind > 'video game'))  
AND ((link_type.id <=> 6  
      OR link_type.id < 6)  
      OR (link_type.link <=> 'referenced in'  
          OR link_type.link < 'referenced in'))  
ORDER BY link_type.id,  
          kind_type.kind,  
          info_type.id  
LIMIT 100;
```

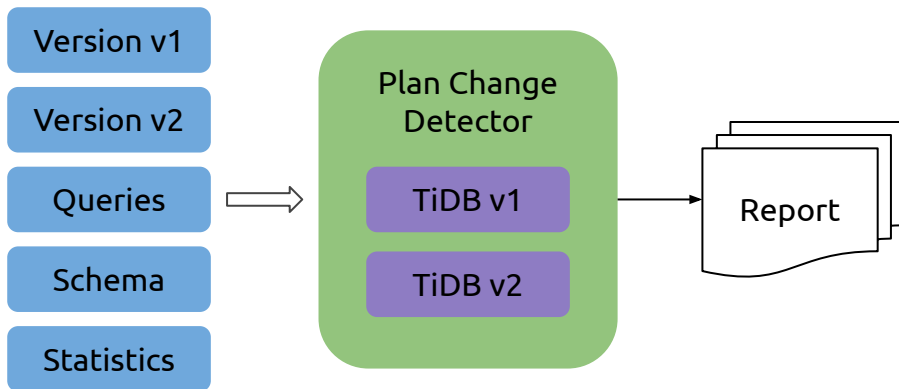
Effectiveness when stats not up-to-date (WIP)

- Split database(e.g IMDB) into slices:
 - Link tables by attribute mapping (primary key \Leftrightarrow foreign key)
 - Organize tables into several groups
 - Split each group by a table or a field
- Incrementally import data slices
- Record the effectiveness metric of each round
- Measure the effectiveness changing




Plan Change Detector (WIP)

Update the cluster
from version v1 to
version v2:



detect plan changes
in each PR:



What's More

Reducing the selectivity estimation error: test the q-error

Learning from the sophisticated DBMS: test their effectiveness



Part III - Optimizer Improvements



Optimizer Improvements

Progress Tracking:

- Root GitHub Issue: [issues/18065](https://github.com/pingcap/tidb/issues/18065)
- Weekly Reports: [Index Selection Work Plan & Weekly Report](#)
- Discuss Here: **#sig-planner**

TiDB Community Slack Channel

<https://pingcap.com/tidbslack/>



Estimation for Out-of-Bound Values (WIP)

Typically when new values are inserted after statistics collected, probably date values ([issues/18461](#))

Old method: $(\text{ModifyRows} / \text{TotalRows}) / \text{NDV}$, unfriendly to small modifications

New method: $1/\text{NDV}$



Avoid Independent Assumptions (WIP)

Extended Statistics ([tidb/issues/18330](https://tidb.io/issue/18330)), similar to PostgreSQL, Oracle

- **create statistics** <stats_name> (<stats_type>) on <tbl> (col [, col])
- **drop statistics** <stats_name>

Supported statistics types:

- cardinality: **where col_a > x and col_b = y**
- correlation: **where col_a > x order by col_b limit 10**



TopN, CM-Sketch and Histogram (WIP)

What are they:

- TopN: Most Common Values (MCV) in PostgreSQL
- CM-Sketch: Two-dimensional Bloomfilter with counters
- Histogram: Ordered Buckets with lower/upper bounds, and counters

Old method:

- TopN is calculated from Histogram ([issues/17467](#)), not removed from Histogram

New method:

- Extract TopN from the data scanned
- Construct CM-SKetch/Histogram without TopN values

Part IV - Join Us!



SQL Engine Team

Improve Optimizer Effectiveness

- Statistics Refactoring
- Index Selection
- Join Order Improvement
- Fast Analyze
- Query Feedback
- Plan Cache
- SQL Plan Management
- Index Tuning Advisor

Build a Fast Query Execution Engine

- Memory Management
- User-Defined Function
- OLTP/OLAP Performance Optimization
- DML Performance Optimization
- Read-Only or Read-Mostly Table Optimization
- Support More Expressions/Executors on Coprocessor

Email: zhangjian@pingcap.com



Quality & Efficiency

Build a Reliabal Database

- Dive into TiDB Implementation, and Try to Destroy
 - TiDB
 - TiFlash
 - Tools
 - DBaas
- New Test Method like Horoscope

Improve Test Efficiency

- Build Effective Test Tools
 - Chaos Mesh®
 - Failpoint
 - Jepsen
- Build an Effective Test Framework, From CI/CD to End to End Test
 - Automation Everything
 - All in K8s

Email: zhangjian@pingcap.com



Thank you!

直播结束后我们会在交流群中分享本期 PPT
还没有加交流群的小伙伴，
可以扫码右边二维码进群👉



发现更多干货文章 & 视频:



微信公众号: PingCAP



知乎专栏: TiDB 的后花园



Bilibili: TiDB_Robot

