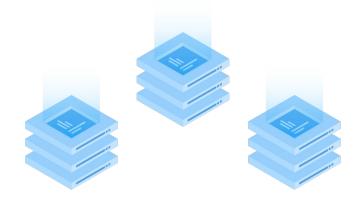
## TiDB SQL 监控及典型的优化案例

qingli.miao











### 课程概要

- 课程背景: TiDB 作为通用分布式数据库产品,和其它数据库一样,也会存在 SQL 性能问题,轻则某个业务 SQL 执行慢,重则某个 SQL就能引发集群整体性能下降或不可用,本课程从原理 + 实践的角度分析导致 slow-query 原因,并且针对常见的原因提供一些优化思路。
- 学习目标:
  - 了解 TiDB 读请求的执行流程。
  - 掌握基于 TiDB 慢查询的故障原因和定位思路。
  - 了解 TiDB 在分析场景下的参数优化。
- 受众:了解和使用过 TiDB 与 MySQL/MariaDB 的学员。
- 时长:1课时(40分钟)
- 关键知识点/大纲:
  - TiDB 读请求的执行流程。
  - slow query 慢查询定位及优化思路。
  - 优化案例。

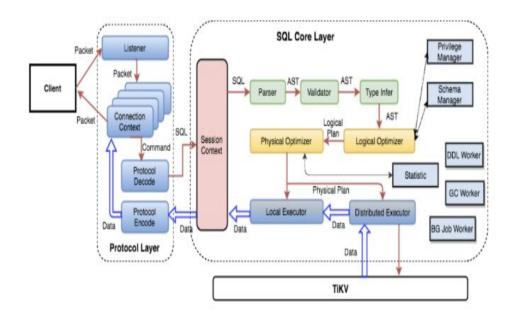




# 读请求执行流程



### SQL 的一生

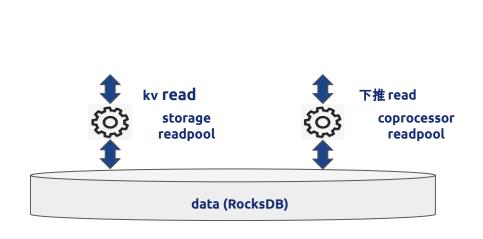






- 从客户端的 Socket 读取一条 SQL
- 获取一个 token
- 从 PD 获取 TSO(事务的时间戳)
- 使用 Parser 将 SQL parse 为 AST
- 将 AST compile 为执行计划
  - Logical Optimizer
  - o Physical Optimizer
  - Execute Plan
- 根据对应的执行计划,最底层的 Executor 会根据这条 SQL 处理的 Key 范围构建出多个要下发到 TiKV 的请求,并通过 distsql 的 API 将这些请求分发到 TiKV
- TiKV 对结果做一些处理(包括 filter, limit等)后,将中间结果集反馈给 TiDB
- TiDB 进行一些表关联、聚合运算等最终反回 给 Client

### SQL 的一生



#### TiKV 收到请求后, 会将请求分为两类:

- storage read pool: 负责主键或唯一索引点查。
- coprocessor:其余的请求。





# 如何定位 slow query



### slow query 产生的原因

- 按组件划分
  - o TiDB
    - parse 慢
    - complie 慢
    - get token 慢
    - 执行计划不正确
  - o PD
    - 获取 tso 慢
  - TiKV
    - 需要扫描大量的 key, 耗时久
    - coprocessor cpu打满,资源等待
    - 读热点

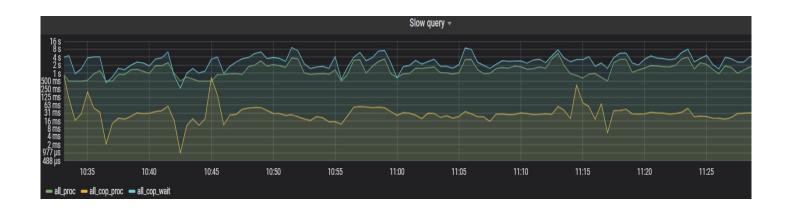






### slow query 的获取渠道

- 集群监控上的 metrics 信息
- slow-query log, 对标 Mysql 格式, 支持市场上的 mysql 慢查询分析工具
- TiDB 的慢查询 SQL 内存表
- TiKV 节点日志 slow-query 反查

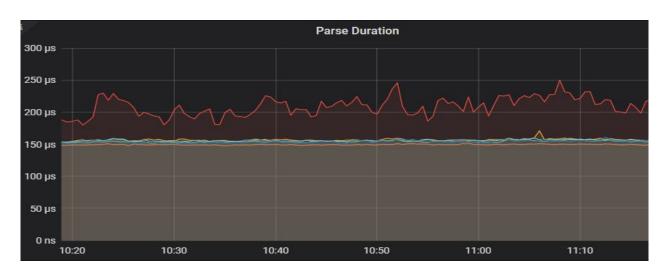






### **TiDB - Parse Metrics**

- 位置: TiDB -> Executor -> Parse Duration
- parse 慢可能原因: TiDB 节点 CPU 压力大
- 图例:

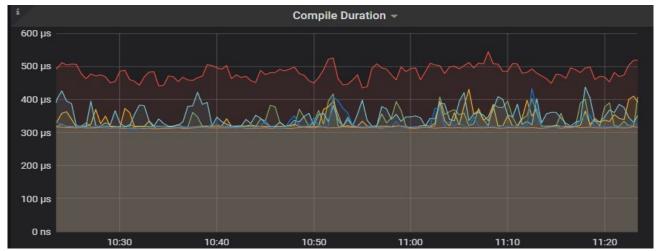






### **TiDB - Compile Metrics**

- 位置: TiDB -> Executor -> Compile Duration
- complie 慢的可能原因:
  - TiDB 节点 cpu 压力大
  - in 子查询结果集多,跟参数 tidb\_opt\_insuqquery\_unfold 有关(2.1)
- 图例:

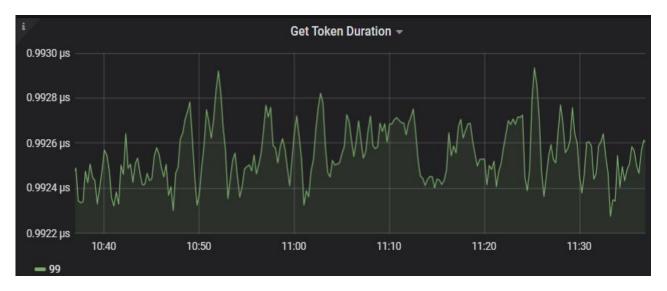






### **TiDB - Get Token Duration Metrics**

- 位置: TiDB -> Server -> Get Token Duration
- 可能慢的原因: token 个数不足, 需要 调整 token-limit
- 图例:

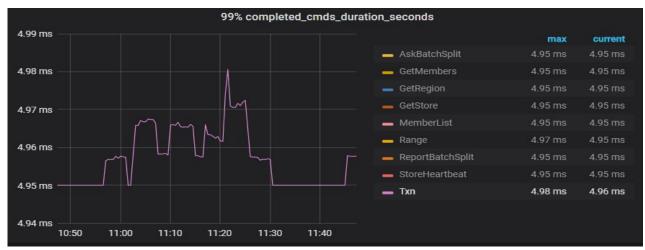






### PD - tso

- 位置: PD -> Grpc -> 99% completed\_cmd\_duration\_seconds -> txn
- 可能慢的原因:
  - PD Leader 切换
  - PD Leader 节点异常,包括 cpu、磁盘等。
- 图例:

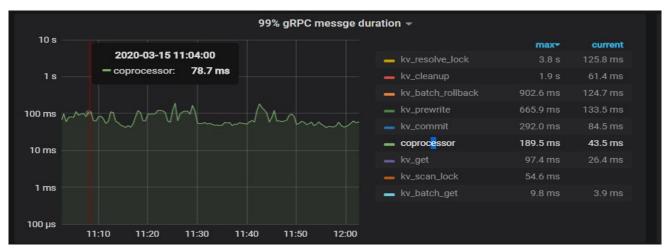






### TiKV - Grpc Duration Metrics

- 位置:TiKV -> Grpc -> Coprocessor
- 可能慢的原因:
  - 需要扫描大量的 key
  - Coprocessor cpu 打满
- 图例:

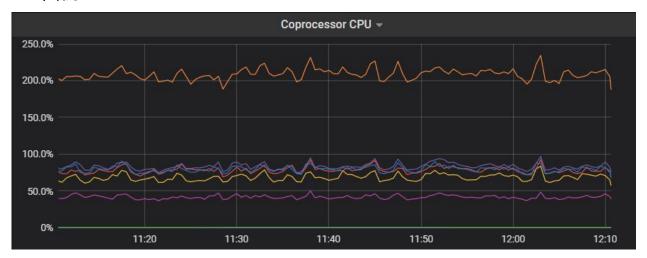






### TiKV - Coprocessor Cpu

- 位置: TiKV -> Thread Cpu -> Coprocessor Cpu
- 可能慢的原因:
  - 扫描大量的 key, 将 cpu 资源占满
  - 读热点, 造成 cpu 资源等待
- 图例:







课程编号: PCTA-104-1

### 慢查询排查技能 - slow query log

```
# Time: 2019-04-25-15:19:33.26029 +0800
# Txn start ts: 407942403346923524
# User: root@127.0.0.1
# Conn ID: 1
# Ouery time: 2.632671582
# Process time: 0.079 Wait time: 0.009 Backoff time: 0.1 Request count: 8 Total keys: 20008 Process keys: 20000
# DB: test
# Index ids: [1]
# Is internal: false
# Digest: edb16a8f28d9c48790925fd1c868fdae3feb49bc58481dda7df228625a5ba6e1
# Stats: t wide:407941920305971202,t slim:pseudo
# Cop_proc_avg: 0.009875 Cop_proc_p90: 0.018 Cop_proc_max: 0.018 Cop_proc addr: 127.0.0.1:22160
# Cop wait avg: 0.001125 Cop wait p90: 0.002 Cop wait max: 0.002 Cop wait addr: 127.0.0.1:24160
# Mem max: 195349
select count(1) from t slim, t wide where t slim.c0>t wide.c0 and t slim.c1>t wide.c1 and t wide.c0 > 5000;
```

- Red color is related with TiDB
- Blue color is related with TiKV/Coprocessor





### 慢查询排查技能 - 内存表

```
TiDB:INFORMATION SCHEMA> select * from `SLOW QUERY` order by
`Query time` desc, Total keys/process keys desc\G
38.854357372
Query time
Process time
          438.314
                       /* /* 有一定的等待时间 */
Wait time
          119.8
                       /* 几乎没有重试. */
Backoff time
            0.104
Request count
           4410
                      /* 旧版本 key 太多 */
Total keys
            2215187529
Process keys
            1108056368
```





TiDB:information schema> show create table `SLOW OUERY` Table Create Table CREATE TABLE `SLOW QUERY` ( SLOW\_QUERY `Time` timestamp unsigned NULL DEFAULT NULL, `Txn\_start\_ts` bigint(20) unsigned DEFAULT NULL, `User` varchar(64) DEFAULT NULL, `Conn\_ID` bigint(20) unsigned DEFAULT NULL, `Ouery time` double unsigned DEFAULT NULL, `Process time` double unsigned DEFAULT NULL, `Wait time` double unsigned DEFAULT NULL, `Backoff\_time` double unsigned DEFAULT NULL, `Request\_count` bigint(20) unsigned DEFAULT NULL, `Total keys` bigint(20) unsigned DEFAULT NULL, `Process\_keys` bigint(20) unsigned DEFAULT NULL, `DB` varchar(64) DEFAULT NULL, `Index\_ids` varchar(100) DEFAULT NULL, `Is internal` tinyint(1) unsigned DEFAULT NULL, `Digest` varchar(64) DEFAULT NULL, `Stats` varchar(512) DEFAULT NULL, `Cop\_proc\_avg` double unsigned DEFAULT NULL, `Cop\_proc\_p90` double unsigned DEFAULT NULL, `Cop proc max` double unsigned DEFAULT NULL, `Cop\_proc\_addr` varchar(64) DEFAULT NULL, `Cop\_wait\_avg` double unsigned DEFAULT NULL, `Cop\_wait\_p90` double unsigned DEFAULT NULL, `Cop wait max` double unsigned DEFAULT NULL, `Cop wait addr` varchar(64) DEFAULT NULL, `Mem\_max` bigint(20) unsigned DEFAULT NULL, `Query` varchar(4096) DEFAULT NULL ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4 bin

### slow log 相关参数

#### slow-threshold

- 参数含义:输出慢 SQL 的耗时阈值, 单位 ms, 静态参数, 需要重启 TiDB SERVER 生效。
- 建议值:在 OLTP 系统, 建议设置 50 ms 左右, OLTP 系统 SQL 一般有以下特点: SQL 短小、单次运行时间短、执行次数多; OLAP 系统可以适当放大点。

#### query-log-max-len

- 参数含义: 日志记录的 SQL 长度, 超过这个长度会截断输出, 单位为字符, 静态参数, 需要重启 TiDB SERVER 生效。。
- 建议值:4096
- tidb\_slow\_query\_file
  - 参数含义:慢查询日志的文件名, 默认值为 tidb\_slow.log, 可以存储 300 MB 内容, 动态参数, 可以通过会话变量 生效。
  - 建议值:tidb\_slow.log 300MB 以后就会重新生成一个文件, 所以如果要定位某一段时间的 slow log, 需要设置该环境变量。TiDB 通过 session 变量 tidb\_slow\_query\_file 控制查询 INFORMATION\_SCHEMA.SLOW\_QUERY 时要读取和解析的文件, 可通过修改 session 变量的值来查询其他慢查询日志文件的内容。





### 用 SQL 从多个维度查询 slow log 表

```
select
query sql text,sum query time,mnt as
executions, avg query time,
(case when avg proc time=0 then 'point get or
commit' when avg proc time>avg wait time then
'coprocessor process' else 'coprocessor wait' end)
as type,
avg proc time, avg wait time, max query time
from (
select substr(query,1,100) query,count(*)
mnt,avg(query time) avg query time
,avg(process time) avg proc time,avg(wait time)
avg wait time, max(query time)
max query time, sum(query time) sum query time
from information schema.slow query
where time >='2020-02-27 05:00:00'
and time <= '2020-02-27 06:00:00'
and lower(query) not like '%analyze%'
or lower(query) not like '%alter%'
group by substr(query,1,100)
) t
order by
sum_query_time desc.
executions desc.
```

avg query time desc

**PingCAP** 

limit 20:

```
PingCAP
University
```

```
text,sum query time,mnt as executions,avq query time, (case when avq proc time=0 then 'point get or commit' when avq proc time>avq wait time then 'coprocessor process' else 'coproces
oc time,avg wait time,max query time from ( select substriquery,1,100) query,count(*) mnt,avg(query time) avg query time ,avg(process time) avg proc time,avg(wait time) avg wait time,max(query time) max query time) max query time,sum(query time) sum c
 ery time from information schema.slow query where time >='2020-82-27 05:00:00' and time <= '2020-82-27 06:00:00' and lower(query) not like '%analyze%' or lower(query) not like '%alter%' group by substr(query,1.100) ) t order by sum qu
 v time desc. executions desc. avo query time desc limit 20:
 sql text
                                                                                                   sum query time
                                                                                                                        executions | avg query time
                                                                                                                                                                              avg proc time
ery_time |
 select 1:
                                                                                                                             3661 | 101.09696439844477 | point get or commit |
                                                                                                    370115.9866627063
                                                                                                                                                                                                                        0 660.
5535779
                                                                                                                            82060 | 0.14186939113911765 | point get or commit |
 commit:
                                                                                                   11641.802236875994
 2642463
 analyze table
                                                                                                    9359.601781304998
                                                                                                                                2 | 4679.800890652499 | point get or commit |
                                                                                                                                                                                                                        0 | 4698.5
 select atr.chlserialno AS chlserialno from
                                                                transaction atr where o.transseri | 7678.476094929006 |
                                                                                                                             19041 | 0.4032601278782105 | coprocessor process |
                                                                                                                                                                                   2.646699186664564 | 1.2161346567932398 | 373.
                Contract set Balance=(select SUM(IFNULL(No.macbalance.0.0)+IFNULL(OverdueBalance.0.0 | 5055.667379390999 |
                                                                                                                              156 | 32.408124226865375 | coprocessor process |
                                                                                                                                                                                  39.130820512820506 | 0.951403846153846 |
                                                                                                                                         1417.277781806 | point get or commit |
 analyze table `arat transa i`;
                                                                                                       2834.555563612
                                                                                                                                                                                                                        0 | 1426.8
 SELECT 0.SERIALNO.0.ACCOUNTNO.0.CONTRACTSERIALNO.0.CUSTOMERID.0.CUSTOMERNAME.0.BUSINESSTYPE.0.PRODUC | 1878.3215532140002 |
                                                                                                                                     4.101138762475983 | coprocessor wait
                                                                                                                                                                               0.0016397379912663763 | 1.7710545851528374 | 185.
analyze table `a
                                                                                                                                2 |
                                                                                                                                          669.988386074 | point get or commit |
                                                                                                       1339.976772148
                                                                                                                                                                                                                        0 670
analyze table `
                                                                                                                                          617.441544562 | point get or commit |
                                                                                                       1234.883089124
                                                                                                                                                                                                                        0 622.2
                                                                                                                              104 | 11.075177008153847 | coprocessor wait
 SELECT 0. SERIALNO.0. PARENTRANSSERIALNO.0. TRANSCODE.0. DOCUMENTTYPE.0. DOCUMENTNO.0. RELATIVEDBJECTTYPE | 1151.8184088480002 |
                                                                                                                                                                                0.025663461538461545
                                                                                                                                     313.0117847546667 | coprocessor wait
 SELECT * FROM cladb Log ONLINE INTERESES t WHERE t.ChannelId = '0
                                                                         3' AND t.update ti
                                                                                                        939.035354264
                                                                                                                                                                                   475.7720000000001 | 2669.1070000000004 |
```

### TiKV 日志反查 slow log

在某个 SQL 执行导致的 TiKV cpu 异常高的场合下, 使用 TiKV 日志可以快速的定位问题 SQL。 TiKV 日志中记录了 slow-query 的几个关键信息:

- ipv4:发出请求的 TiDB 地址, 通过该地址可以确定 SQL 在哪个 TiDB 节点执行。
- start\_ts:事务的 start\_ts, 通过 start\_ts 反查 slow log, 快速定位问题 SQL。
- table\_id:查询的那张表,可以通过 information\_schema.tables 的 tidb\_table\_id 反查表名

```
[2020/03/13 14:23:32.605 +08:00] [INFO] [raft.rs:723] ["[region 1097830] 1823912 [logterm: 26, index: 53, vote: 0] cast MsgRequestVote for 1772152 [logterm: 26, index: 53] at term 27"]
[2020/03/13 14:23:32.605 +08:00] [INFO] [tracker.rs:175] ["[region 1097830] 1823912 [logterm: 26, index: 53, vote: 0] cast MsgRequestVote for 1772152 [logterm: 26, index: 53] at term 27"]
[2020/03/13 14:23:32.605 +08:00] [INFO] [tracker.rs:175] ["[region 1097830] 1823912 [logterm: 26, index: 53, vote: 0] cast MsgRequestVote for 1772152 [logterm: 26, index: 53] at term 27"]
[2020/03/13 14:23:32.605 +08:00] [INFO] [tracker.rs:175] ["[region 1097830] 1823912 [logterm: 26, index: 53, vote: 0] cast MsgRequestVote for 1772152 [logterm: 26, index: 53] at term 27"]
[2020/03/13 14:23:32.605 +08:00] [INFO] [tracker.rs:175] ["[region 10978090] [370-query] execute takes 4.78s, wait takes 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(415257214799839333), table_id: Some(1391), tag: select takes: 17ms, peer: Some("iput" "43041"), start.ts: Some(41525721479983933
```





### 使用方法总结

- 通过监控的相关 metrics 可以获取到 TiDB 集群的整体运行情况, 比如某些 TiKV 节点 CPU 高不高等。
   对于定位慢查询来说可以有一个初步的认知, 用 metrics 来判断那些时间段存在问题, 可能存在那些问题。
- slow query log 及 slow log 内存表,可以从整体上观察 SQL 运行情况,比如 SQL 执行花了多长时间, 执行了多少次,哪些 SQL 最耗时,哪些 SQL 占用资源最多等等。
- 但是 slow log 内存表还存在一定的局限性,因 为每个 TiDB 有自己的一个 slow log 且 slow log file 到 达 300 mb 后还会切换,所以使用起来还是有一定的局限性。
- TiKV 日志能够快速定位问题 SQL,对于只执行一次或几次的大 SQL 来说,非常容易定位问题。但是如果是单次执行快但是执行频率高的 SQL 来说,查看 TiKV 日志还是相对比较麻烦。
- 综上所述,结合的使用三个工具才能比 较快速定位问题 SQL。







# AP 场景加速



### 相关并发参数

#### tidb\_distsql\_scan\_concurrency

这个变量用来设置 scan 操作的并发度,对于 AP 类应用, 最大值建议不要超过所有 TiKV 节, 的 CPU 核数

#### tidb\_index\_serial\_scan\_concurrency

这个变量用来设置顺序 scan 操作的并发度

#### tidb\_index\_lookup\_concurrency

这个变量用来设置 index lookup 操作的并发度

#### tidb\_index\_lookup\_join\_concurrency

这个变量用来设置 index lookup join 算法的并发度

#### tidb\_hash\_join\_concurrency

这个变量用来设置 hash join 算法的并发度





# SQL 优化案例



### TiDB SQL 优化需要注意的点

- where + col1 条件 or col2 条件,在 TiDB 中无法使用多个索引即 MySQL 的 index\_merge 功能,会导致全表扫描,可以转换为 from xxx where col1 union from xxx where col2,这样就可以同时使用 col1索引和 col2索引。注意:需要注意本身结果集是否有重复数据,若有重复数据 union 和 or 无法等价转换。
- index join 的被驱动表,不能有函数处理的过滤条件,类似 where year(xxx)='2019',尽量不要在过滤条件增加对列的函数处理。(现在还不支持函数索引)
- 其余与通用数据库 SQL 一样,需要注意尽量减少 not in, not exits 使用,尽量用表关联。
- 亮点: TiDB 支持 null 及 like 'abc%' 使用索引, !=xxx 在 TiDB 也能转换为索引范围扫描, 可以转换为 [-inf,xxx),(xxx,+inf], 在某些特定场景比较有意义。







### Union + 函数改写案例

#### sql 文本:

```
select r.*,m.real steel grade,m.house code id,m.read meter type from r.d. actor m inner join
p.meter plan code id,read contract code id,meter code id,meter plan no,use gas type,company code no,meter unit
code, meter type, meter read reason,
meter_reading_emp_name,meter_plan_status,plan_meter_reading_date,last_actual_meter_date,actual_read_meter_date,
meter_round_count,meter_residual_count,meter_used_count
    .0.0000 last reading count,0.0000 current meter count,0.0000 current gas count,0.00 payable money
    from and_mater_plan_code_id =
c.meter_plan_code_id
    union
    select
p.meter_plan_code_id,read_contract_code_id,meter_code_id,meter_plan_no,use_gas_type,company_code_no,meter_unit_
code,meter_type,meter_read_reason,
meter reading emp name, meter plan status, plan meter reading date, last actual meter date, actual read meter date,
    0.0000 meter round count, 0.0000 meter residual count, 0.0000 meter used count,
    last_reading_count,current_meter_count,current_gas_count,payable_money
    from and meter plan record into prinner join and meter column, record detail c on p.meter plan code id =
c.meter_plan_code_id
    on m.meter_code_id=r.meter_code_id where m.house code id in
    ('C000-4454be745
       and year(actual read meter date) = '2019'
```



order by house\_code\_id desc, meter\_code\_id desc



#### 执行计划:

id	count	task	operator info	execution info
a Sort 20	0.00		operator into eceigas master in house code ididesi; rimeter code ididesi;	time:3m55.74293361s, loops:2, rows:1
W15-01	-	root	eceggasmaster.m.house_code_jd.cess, r.meter_code_jd.dess. r.meter_plan code id, r.read contract code id, r.meter code id, r.meter plan no, r.use gas type, r.company code no, r.meter urit code, r.meter	Transfer and the second
L-Projection_22	0.00	1000		
L-HashRightJoin_29		root	inner join, inner Indexi.cokUp_85, equal(eq(ecejgasmaster.m.meter_code_jd, r.meter_code_jd))	time:3m55.742868135s, loops:2, rows:1
HindexLookUp_85	0.00	root		times1.739923ms, loops:2, rows:2
I ⊢IndexScan_83	0.00	сор	tablem, indexhouse_code_id, rangel[10000c4451be748d0b93278e46719c396",10000c4451be748d0b93278e46719c336"], keep ordenfalse	
LableScan_84	0.00	сор	tablemd_meter, keep ordenfake	
-HashAgg_32	33620595.80	root	group byzactual read meter date, company code no, current gas count, current meter count, last actual meter date, last reading count, met	
└─Urion_33	33620595.80	root		time:1m26.777315355s, loops:929547, rous:29745388
I-Projection_34	69695.00	root	ecojgasmaster,p.meter plan code id, ecojgasmaster,p.read contract code id, ecojgasmaster,p.meter code id, ecojgasmaster,p.meter plan no	
\_Projection_35	69695.00	root	ecojgasmaster.p.meter_plan_code_id, ecojgasmaster.p.read_contract_code_id, ecojgasmaster.p.meter_code_id, ecojgasmaster.p.meter_plan_nc	
└─HashLeftJoin_41	69695.00	root	inner join, innenTableReader 56, equal(eq(ecejgasmaster,p.meter_plan_code_id, ecejgasmaster.c.meter_plan_code_id)]	time:40.667398452s, loops:735, rows:23410
I ├─Selection_52	33550900.80	root	eq(year(ecejgasmaster,p.actual_read_meter_date), 2019)	time:39.465949592s, loops:929004, rows:29728065
└─TableReader_54	41938626.00	root	data:TableScan_53	time:27.043736101s, loops:1310578, rows:41938432
1   _TableScan_53	41938626.00	cop	table:p, rangetj-inf,+inf), keep ordenfalse	
1 L—TableReader_56	69695.00	root	data:TableScan_55	time:119.359342ms, loops:2180, rows:69704
LableScan_55	69695.00	cop	tablex; ranged-inf,+inf), keep orderfalse	
L-Projection_57	33550900.80	root	ecejgasmaster.p.meter_plan_code_id, ecejgasmaster.p.meter_plan_node_id, ecejgasmaster.p.meter_code_id, ecejgasmaster.p.meter_plan_node.	, ec time:1m6.700690075s, loops:928813, rows:29721976
L-Projection_58	33550900.80	root	ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.mead_contract_code_jd, ecejgasmaster.p.meter_code_jd, ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.meter_code_jd, ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.meter_plan_code_jd, ecejgasmaster.p.meter_code_jd, ecejgasmaster.p.met	, ec time:1m31.154399383s, loops:928813, rows:29721976
⊢Mergeloin_59	33550900.80	root	inner join, left keyecejgasmaster.p.meter plan code id, right keyecejgasmaster.cmeter plan code id	time:1m33.569924265s, loops:928813, rows:29721976
I—Selection_66	33550900.80	root	eq(year(ecejgasmaster.p.actual_read_meter_date), 2019)	time:34.386847091s, loops:929005, rows:29728065
I └─IndexLockUp_70	41938626.00	root		times22.438786601s, loops:1310579, rows:41938432
I ⊢IndexScan_68	41938626.00	cop	tableip, indexmeter_plan_code_id, range([NULL+inf), keep ordentrue	
└─TableScan_69	41938626.00	сор	tablemd_meter_plan_record_info, keep ordenfalse	
└─IndexLookUp_74	41947459.00	root		time:25.153792414s, loops:1310796, rows:41945411
├─IndexScan 72	41947459.00	cop	tablex, index.meter plan code id, range:(NULL, +inf), beep order true	
L-TableScan 73	41947459.00	сор	table mid meter ordinary record detail, keep orderfalse	

### Union + 函数改写案例

#### 改写后sql 文本:

```
select t.* from(
select xxx,m.read meter type
from p inner join m on m.meter code id=p.meter code id
inner join md meter card record detail c on p.meter plan code id
= c.meter plan code id where
m.house code id in ('xxx')
union
select xxx,m.read meter type
from p inner join m on m.meter code id=p.meter code id
inner join md meter ordinary record detail c on
p.meter plan code id = c.meter plan code id
where
m.house code id in (xxxx)
where
t.actual read meter date >= '2019-01-01 00:00:00'
and t.actual read meter date < '2020-01-01 00:00:00'
order by
t.house code id desc.
t.meter plan code id desc
```

#### 改写后执行计划:

t.house code id:desc, t.meter plan code id:desc

			· ·
Sort	_22	2.00	root
L-H	ashAgg_27	2.00	root
L_	Union_28	0.00	root
H	-Projection_29	0.00	root
- 1	L—Projection_30	0.00	root
- 1	L—IndexJoin_35	0.00	root
- 1	-IndexJoin_67	0.00	root
1	IndexLookUp_66	12.13	root
- 1	-IndexScan_63	16.95	сор
- 1	L—Selection_65	12.13	сор
-1	I I └─TableScan_6	16.95	сор
-1	└─IndexLookUp_85	0.00	root
-1	IndexScan_83	0.00	сор
1	I └─TableScan_84	0.00	сор
- 1	└─IndexLookUp_34	1.00	root
1	IndexScan_32	1.00	сор
- 1	L—TableScan_33	1.00	сор
L	-Projection_91	0.00	root
	L-Projection_92	0.00	root
	└─IndexJoin_97	0.00	root
	IndexJoin_129	0.00	root
	-IndexLookUp_12	12.13	root
	-IndexScan_125	16.95	сор
	L—Selection_127	12.13	сор

group by:actual\_read\_meter\_date, company\_code\_no, current\_gas\_count, current\_m
ecejgasmaster.p.meter\_plan\_code\_id, ecejgasmaster.p.read\_contract\_code\_id, ecejg

ecejgasmaster.p.meter\_plan\_code\_id, ecejgasmaster.p.read\_contract\_code\_id, ecejg inner join, inner:IndexLookUp\_34, outer key:ecejgasmaster.p.meter\_plan\_code\_id, ir inner join, inner:IndexLookUp\_66, outer key:ecejgasmaster.m.meter\_code\_id, inner l

table:p, index:meter\_code\_id, range: decided by [ecejgasmaster.m.meter\_code\_id], ge(ecejgasmaster.p.actual\_read\_meter\_date, 2019-01-01 00:00:00.000000), lt(ecejga table:md meter plan record info, keep order:false

table:m, index:house\_code\_id, range:["0000c4451be748d0b93278e46719c336","000 table:md\_meter, keep order:false

table:c, index:meter\_plan\_code\_id, range: decided by [ecejgasmaster.p.meter\_plan\_table:md\_meter\_card\_record\_detail, keep order:false ecejgasmaster.p.meter\_plan\_code\_id, ecejgasmaster.p.meter\_plan\_code\_id, ecejgasmaster.p.meter\_plan\_code\_id, ecejgasmaster.p.meter\_plan\_code\_id, ecejgasmaster.p.meter\_plan\_code\_id, ir

table:p, index:meter\_code\_id, range: decided by [ecejgasmaster.m.meter\_code\_id], qe(ecejgasmaster.p.actual\_read\_meter\_date, 2019-01-01 00:00:00.000000), [t(ecejga

inner join, inner:IndexLookUp\_128, outer key:ecejgasmaster.m.meter\_code\_id, inner











### 相关子查询优化

#### sql 文本:

LEFT JOIN xxx mord ON mord.meter\_plan\_code\_id = plan.meter\_plan\_code\_id - WHERE c.customer\_code\_id = 'xxxx' AND h.house\_code\_id = 'xxxx' AND m.meter\_code\_id= 'xxxxx' AND mmeter\_code\_id= 'xxxxx' AND mpr.create\_time = (
SELECT MAX(create\_time) FROM xxx price WHERE price.read\_contract\_code\_id = mrc.read\_contract\_code\_id = mrc.read\_contract\_code\_id = (
SELECT mpri.meter\_plan\_code\_id FROM xxxxmpri WHERE mpri.meter\_code\_id = m.meter\_code\_id ORDER BY create\_time DESC LIMIT 1

#### 执行计划:

id		count	task	operator into		
11	LIndexRead	d 1.00	root	index:IndexScan_87		
11	└─IndexSca	1.00	сор	table:rcp, index:read_contract_code_id, range: decided by [ecejgasmaster.mrc.read_contract_code_id], keep ord		
11	L—IndexLookU	r 1.00	root			
11				table:mpr, index:read_contract_code_id, range: decided by [ecejgasmaster.mrc.read_contract_code_id], keep order		
11				table:md_price_record, keep order:false		
1.1	LIndexLookUp_20.00 root					
11		0.00	сор	table:plan, index:meter_code_id, range:["290905a5b7b544cba844d54fcfad4d4d","290905a5b7b544cba84d54fcfad4d4d","290905a5b7b544cba84d54fcfad4d4d","290905a5b7b54dcba84d54fcfad4d4d","290905a5b7b54dcba84d6d5dfcfad4d4d","290905a5b7b54dcba84d6dfcfad4dfdfcfad4dfdfcfad4dfdfdfdfdfdfdfdfdfdfdfdfdfdfdfdfdfdf		
11	LableScan_20	0.00	сор	table:md_meter_plan_record_info, keep order:false		
11	III LableReader_206	5 2.64	root	data:Selection_205		
11	III L—Selection_205	2.64	сор	eq(ecejgasmaster.dm.del_flag, 0), eq(ecejgasmaster.dm.dict_type_key, "meter_location")		
11	L—TableScan_204	4 326.00	сор	table:dm, range:[-inf,+inf], keep order:false		
11		1.00	root	index:IndexScan_74		
11	I I LIndexScan_74	1.00	сор	table:card, index:meter_plan_code_id, range: decided by [ecejgasmaster.plan.meter_plan_code_id], keep order:f		
11	└─IndexReader_66	1.00	root	index:IndexScan_65		
11	└─IndexScan_65	1.00	сор	table:mord, index:meter_plan_code_id, range: decided by [ecejgasmaster.plan.meter_plan_code_id], keep order		
1	└─StreamAgg_215	2676401.7	root	group by:ecejgasmaster.price.read_contract_code_id, funcs:max(ecejgasmaster.price.create_time), firstrow(ecejgasmaster.price.create_time)		
1	L—IndexLookUp_226	2731537.0 root				
1	IndexScan_224	2731537.0 cop		table:price, index:read_contract_code_id, range:[NULL,+inf], keep order:true		
1	L—TableScan_225	2731537.0	сор	table:md_price_record, keep order:false		





### 相关子查询优化

#### 改写后 sql 文本:

#### 改写后的执行计划:

LEFT JOIN xxx mord ON mord.meter plan code id = plan.meter plan code id inner join(SELECT MAX(create time) as max create time, read contract code id FROM ( SELECT price.create time, price.read contract code id FROM xxxx price JOIN(SELECT mrc.read contract code id from xxxx AS h left JOIN xxxx AS m ON m.house code id = h.house code id left JOIN xxx AS mrc ON mrc.meter code id = m.meter code id where h.house code id = 'xxxx' AND m.meter code id = 'xxx') tmp tWHERE price.read contract code id = tmp t.read contract code id) price GROUP BY price.read contract code id) price ON mpr.create time = price.max create time and price.read contract code id = mrc.read contract code id WHFRF c.customer code id = 'xxx' AND h.house code id = 'xxxx' AND m.meter code id = 'xxxx'

1	I IndexScan_74	1.00	сор	table:mord, index:meter_plan_code_id, range: decided by [ecejgasmaster.plan.meter_plan_code_id], keep order:f
1	L—HashAgg_221	1.00	root	$group\ by: ecejgas master. price. read\_contract\_code\_id,\ funcs: max(ecejgas master. price. create\_time),\ firstrow(ecejgas master. price. create\_time),\ firstrow(e$
I	L-IndexJoin_227	0.00	root	$inner\ join, inner\ Index Look Up\_226, outer\ key: ecej gas master. mrc. read\_contract\_code\_id, inner\ key: ecej gas master.$
Î	⊢IndexLookUp_226	1.00	root	
I	IndexScan_224	1.00	сор	table:price, index:read_contract_code_id, range: decided by [ecejgasmaster.mrc.read_contract_code_id], keep or
ĺ	LableScan_225	1.00	сор	table:md_price_record, keep ordentalse
I	L—IndexJoin_234	0.00	root	left outer join, inner:IndexLookUp_233, outer key:ecejgasmaster.m.meter_code_id, inner key:ecejgasmaster.mrc.n
١	HashLeftJoin_236	0.00	root	inner join, inner:IndexLookUp_243
1	IndexReader_23 142.53		root	index:IndexScan_238
I	L—IndexScan_23{142.53		сор	table:h, index:house_code_id, range:[*1790fdf0c2de42ca955a5c45e788036e*,*1790fdf0c2de42ca955a5c45e788031
١	I LIndexLookUp_240.00		root	
İ	I I I—IndexScan 240 1 00		con	tablem indermeter code id ranne/1°290905a5h7h544cha844d54fcfad4ddd" "290905a5h7h544cha844d54fcfad4a





# 谢谢!







### 课程作业

- 问题1:slow query 内存表( )字段代表着 SQL 实际的查询时间。 A. query\_time B.process\_time C.wait\_time D.execute\_time
- 问题2:点查请求在 TiKV 中是通过()执行的。 A. coprocessor B.key range C.storage read pool D.parse
- 问题3:在 TiDB中 以下哪个条件无法走索引?() A.like 'abc%' B.!= C.in D.or
- 问题4: Get Token duration 是指?()
  A. SQL 解析时间 B. 获取 PD tso 时间 C.获取 token 的时间 D.生成执行计划时间
- 问题5:某几个 TiKV 节点在 10:00 11:00 CPU 使用率在 90%以上, 如何定位是哪个或那些 SQL 影响到了?





